

Relative Wealth Concerns, Network Aggregation, and Effective Risk Aversion

Renato Vicente*

August 22, 2026

Abstract

In a one-period CARA economy with Gaussian payoffs, exogenous quantities and endogenous prices, each investor compares his wealth to a weighted average over a network of peers. Equilibrium demand scales solve a Katz–Bonacich system, and prices are explicit in one scalar, the effective market coefficient, which is itself independent of prices; the market Sharpe ratio is proportional to it. Two limits, available for every network, organise how the network aggregates risk aversion. At zero comparison intensity the marginal effect is a quadratic form in risk tolerances, and a doubly stochastic influence matrix — balanced attention — makes it a Dirichlet form, hence non-negative for every profile of tolerances. Under intense comparison the coefficient converges to a stationary average of individual coefficients, weighted by influence rather than by population. It then falls below its no-comparison value, taking the Sharpe ratio with it, whenever stationary influence sufficiently favours the less risk averse; it interpolates from the harmonic to the arithmetic mean when influence is spread evenly; and it does not move at all when the network splits into blocks homogeneous in type. A separate result is that comparison is exactly neutral, for every network and every profile of intensities, when net risk aversion is common.

Keywords: relative wealth concerns, keeping up with the Joneses, CARA aggregation, Katz–Bonacich centrality, homophily.

JEL: G11, G12, D53, D85.

1 Introduction

[12] studies an exchange economy in which agents evaluate consumption relative to a weighted average of the consumption of others, and asks when equilibrium prices depend

*Dept. of Applied Mathematics, Instituto de Matemática, Estatística e Ciência da Computação, Universidade de São Paulo. `rvicente@ime.usp.br`

only on aggregates. His Proposition 1 shows that linear risk tolerance with common cautiousness delivers a representative agent and rules out portfolio bias; his Proposition 2 gives the conditions under which prices depend on the aggregate endowment alone, and finds them restrictive: the preference parameters must be common across agents and, absent the trivial case of uniform weights, so must the equilibrium value of the endowment. Under those conditions the CARA case reduces cleanly, and keeping up with the Joneses is equivalent to lowering the representative agent's risk aversion from α to $\alpha(1 - \gamma)$.

What happens outside those conditions is the natural question, and [12] leaves it open. He notes that in their absence equilibrium prices may be affected by the distribution of weights in the definition of the Joneses, by the preference parameters, and by the endowment distribution — three channels that a representative-agent reduction cannot express. This paper opens the first of them, in the CARA case, by letting the benchmark be agent-specific: investor a compares himself to $\sum_j A_{aj}W_j$, a weighted average over a network of influence rather than a single economy-wide aggregate. The question is then not whether the comparison matters but *what feature of the network* decides how it matters.

Two remarks on what is being held fixed, since they govern the reading of everything below. First, with preferences $-\exp[-\theta_a W_a - \lambda_a(W_a - \sum_j A_{aj}W_j)]$, the coefficient of absolute risk aversion in own wealth is $\alpha_a = \theta_a + \lambda_a$, not θ_a . We call θ_a *net* risk aversion: it is the coefficient on a common translation of all wealths, and it is risk aversion net of the comparison. Raising λ_a at fixed θ_a therefore raises own-wealth risk aversion too, and the comparative statics below are taken along that path. Second, the economy is closed: payoffs and quantities are the primitives, prices are endogenous, and Section 3.2 shows that equilibrium prices are explicit in the primitives and in a single scalar $\bar{\theta}$ which is itself independent of prices. The market Sharpe ratio equals $\bar{\theta}$ times a constant on the payoff side, so the object characterised below is a price of risk in the strict sense.

The demand scales solve $[\mathbb{D}_{\theta+\lambda} - \mathbb{D}_\lambda A]\vec{c} = \mathbf{1}$, a Katz–Bonacich system in the sense of [2], and the paper's contribution is to characterise how that system aggregates. Two limits, both available for every network, do the work.

At the origin, the marginal effect of comparison is the quadratic form $\bar{\theta}'(0) = m \vec{t}^\top (\mathbb{I} - A) \vec{t} / (\mathbf{1}^\top \vec{t})^2$ in the vector $\vec{t}$ of risk tolerances (Theorem 2). If attention is balanced — A doubly stochastic, so that each investor receives as much attention as he gives, and not necessarily symmetric — this is a Dirichlet form, $\frac{1}{2} \sum_{a,j} A_{aj} (t_a - t_j)^2$, hence non-negative and zero only if tolerance is constant along every link. Comparison raises the effective coefficient exactly to the extent that investors look at people unlike themselves.

At the other end, as comparison becomes intense, $\bar{\theta}$ converges to a *stationary* average of individual coefficients: $\vec{\pi}^\top \vec{\theta}$ when the influence matrix has a single recurrent class, and a harmonic combination across classes in general (Theorem 3). Influence, not population, does the weighting in the limit, and since the no-comparison value is the harmonic mean $H(\theta)$, a sufficient condition for comparison to *lower* the effective coefficient is $\vec{\pi}^\top \vec{\theta} < H(\theta)$.

These two limits organise what would otherwise be a list. The mean-field regime

is the case in which the stationary distribution is uniform, so the limit is $\mathbb{E}[\theta]$ and comparison interpolates from the harmonic to the arithmetic mean (Proposition 1). Sign reversal is the case in which stationary influence favours the less risk averse; in the three-investor example of Proposition 2 the most risk averse investor is watched by nobody, receives stationary weight zero, and the limit falls to $3/4$ against a starting value of $6/7$. Perfect homophily is the case of several recurrent classes each homogeneous in type, where the stationary average reproduces the harmonic mean and the two limits coincide, so the aggregate never moves at all (Proposition 3).

A separate and much simpler result delimits where any of this can happen. If net risk aversion is common, the comparison is exactly neutral: every investor holds the position he would hold with no relative concern, for every row-stochastic A and every profile λ_a (Theorem 1). No statistic of the network survives into equilibrium. This is the network extension of Gómez’s CARA case, and it says that dispersion of net risk aversion is the necessary fuel for everything else.

The message is that the comparison channel in this class of models is a property neither of preferences alone nor of the network alone, but of the alignment between them — specifically, of the correlation between net risk aversion and stationary influence. A representative-agent treatment delivers the mean-field case, and can therefore be wrong in sign, not merely in magnitude, whenever attention is unbalanced.

2 Related literature

Preferences over relative consumption or wealth have been used to address a range of asset-pricing questions. [1] and [10] build consumption-based models with external benchmarks and obtain an amplified equity premium; [3] show how an external habit generates time-varying effective risk aversion; [4] make the wealth distribution a state variable under heterogeneous CRRA preferences, which is the dynamic counterpart of the configuration studied here. [9] provide the taxonomy under which (2) is jealousy combined with keeping up with the Joneses.

Within the CARA class, [5] is the closest antecedent and the comparison locates what is new here. They study two agents with multiple risky assets and distributions in a natural exponential family, and show that relative concerns act through a transformation of effective risk tolerances, lowering equilibrium premia. The mechanism they identify is a change in the tolerance *levels*. With two agents the benchmark of each is the other, so there is no network to speak of: the influence matrix is forced to $A = \begin{pmatrix} 0 & 1 \\ 1 & 0 \end{pmatrix}$, which is doubly stochastic, and by Theorem 2 the marginal effect is then necessarily non-negative. The questions of this paper — whether attention is balanced, which investors stationary influence favours, whether types and positions are correlated — cannot be posed at $m = 2$, and the results that turn on them (Theorems 2 and 3, and the reversal and homophily that follow) have no two-agent counterpart. What is added is therefore not a dimension count but the object doing the aggregating: a stationary distribution in place of a pair of tolerances. [8] obtain portfolio bias when the benchmark differs in composition from the investor’s own holdings, a channel excluded here by the scalar

structure of (6). [14] solve relative-performance games for CARA and CRRA investors in mean-field and n -agent form; their concern is the existence and structure of Nash equilibria rather than whether the aggregate moves, and their mean-field formulation is the analogue of our Proposition 1 in a different setting. [13] take the problem to incomplete markets in continuous time. [16] surveys what heterogeneity and inequality imply for asset prices more broadly.

Two connections deserve care. [11] gives the criterion governing the effect of endowment inequality on the equity premium, namely concavity of risk tolerance in wealth. Our (21) is an analogue for dispersion in a preference parameter rather than in endowments, and should be read as such rather than as an application of his theorem. And [2] give the linear-quadratic network game whose solution (7) reproduces; what the present setting adds is that the centrality vector becomes constant, and every statistic of the network drops out, precisely when net risk aversion is common.

3 The economy

The primitives are payoffs and quantities, not return moments. There are two dates. At the second date n risky assets pay $\vec{X} \in \mathbb{R}^n$, Gaussian with mean $\vec{\mu}$ and positive definite covariance V ; they are in exogenous supply of $\vec{\nu} > 0$ shares and trade at the first date at prices $\vec{p} \in \mathbb{R}^n$, to be determined. A riskless asset pays gross return $R_f > 0$. There are m investors; investor a begins with wealth W_a^0 and chooses share holdings $\vec{z}_a \in \mathbb{R}^n$, placing the remainder in the riskless asset, so that

$$W_a = R_f(W_a^0 - \vec{p}^\top \vec{z}_a) + \vec{X}^\top \vec{z}_a. \quad (1)$$

Preferences are

$$U_a = -\exp \left[-\theta_a W_a - \lambda_a \left(W_a - \sum_j A_{aj} W_j \right) \right], \quad \theta_a > 0, \lambda_a \geq 0. \quad (2)$$

Assumption 1 (Normalised benchmark). $m \geq 2$; $A_{aa} = 0$, $A_{aj} \geq 0$ and $\sum_j A_{aj} = 1$ for every a .

Definition 1 (Two coefficients). Holding the wealth of others fixed, the coefficient of absolute risk aversion of (2) in own wealth is

$$\alpha_a = \theta_a + \lambda_a, \quad (3)$$

which we call *own-wealth risk aversion*. We call θ_a the *net risk-aversion coefficient*: under Assumption 1 it is the coefficient on a common translation of all wealths, since adding a constant to every W_j shifts the exponent of (2) by $-\theta_a$ times that constant, independently of λ_a . It is risk aversion net of the comparison.

The distinction matters for reading everything below, and we state it before any result because the two coefficients cannot both be held fixed. Raising λ_a at constant θ_a

is not a pure increase in the intensity of comparison: by (3) it raises own-wealth risk aversion by the same amount. Every comparative static in λ in this paper is taken along that path, at constant θ , and the neutrality of Section 4 is exactly the cancellation of the two effects. The alternative path, at constant α , is described in Remark 2.

The multiplicative parametrization of [12] makes the same point in his notation. Collecting terms in the exponent of (2),

$$-\theta_a W_a - \lambda_a \left(W_a - \sum_j A_{aj} W_j \right) = -\alpha_a \left(W_a - \gamma_a \sum_j A_{aj} W_j \right), \quad \gamma_a = \frac{\lambda_a}{\theta_a + \lambda_a} \in [0, 1), \quad (4)$$

so that $\theta_a = \alpha_a(1 - \gamma_a)$. He writes the CARA case as $u(c, X) = -\exp(-\alpha(c - \gamma X))$ and reduces it to a representative agent with coefficient $\alpha(1 - \gamma)$, which is exactly θ .

The benchmark is thus a weighted average of the wealth of the other investors, as in [12], with the difference that the weights may vary across investors. Assumption 1 is what gives θ_a the reading in Definition 1, and it is what the results below require. Dropping it changes the object being compared: row sums below one describe comparison with a partial or external benchmark, which is a coherent alternative we do not study, while row sums above $1 + \theta_a/\lambda_a$ make an investor worse off when the whole economy becomes richer, which is harder to motivate. We retain Assumption 1 throughout and claim nothing about either case.

Assumption 2 (Nash behaviour). Each investor chooses $\vec{z}_a$ taking the holdings $\vec{z}_j$, $j \neq a$, as fixed at their equilibrium values. An equilibrium is a profile $(\vec{z}_1, \dots, \vec{z}_m)$ and a price vector $\vec{p}$ such that each $\vec{z}_a$ is optimal given the others and markets clear.

Assumption 2 is the standard notion and we state it explicitly because the conjecture it embodies matters for what follows. An investor behaving in this way conjectures that the positions of others do not respond to his own: the conjectural derivative is zero, and the conjectured levels are the equilibrium positions.

3.1 First-order conditions and equilibrium prices

Substituting (1) into (2), the random part of the exponent is $-\vec{X}^\top \vec{q}_a$ with

$$\vec{q}_a = (\theta_a + \lambda_a) \vec{z}_a - \lambda_a \sum_j A_{aj} \vec{z}_j, \quad (5)$$

and the deterministic part is $-(\theta_a + \lambda_a) R_f(W_a^0 - \vec{p}^\top \vec{z}_a) + \lambda_a \sum_j A_{aj} R_f(W_j^0 - \vec{p}^\top \vec{z}_j)$. Gaussian payoffs give $\mathbb{E}[U_a] = -\exp(d_a - \vec{q}_a^\top \vec{\mu} + \frac{1}{2} \vec{q}_a^\top V \vec{q}_a)$ with d_a the deterministic part. Since $\partial \vec{q}_a / \partial \vec{z}_a = (\theta_a + \lambda_a) \mathbb{I}$, the exponent is a strictly convex quadratic in $\vec{z}_a$, with Hessian $(\theta_a + \lambda_a)^2 V$. The optimum is therefore interior and solves

$$V \vec{q}_a = \vec{\mu} - R_f \vec{p}. \quad (6)$$

Every solution is proportional to $V^{-1}(\vec{\mu} - R_f \vec{p})$, so all investors hold the same portfolio of risky assets and differ only in scale; two-fund separation holds, as in Gómez's Proposition 1. Writing $\vec{z}_a = c_a V^{-1}(\vec{\mu} - R_f \vec{p})$ reduces (6) to

$$[\mathbb{D}_{\theta+\lambda} - \mathbb{D}_\lambda A] \vec{c} = \mathbf{1}, \quad \mathbb{D}_x = \text{diag}(x_1, \dots, x_m), \quad (7)$$

whose solution is a Katz–Bonacich centrality of the influence network [2], generalised in that the attenuation factors λ_a and the base vector are agent-specific rather than scalar. Write $M = \mathbb{D}_{\theta+\lambda} - \mathbb{D}_\lambda A$. Its off-diagonal entries are non-positive, so M is a Z-matrix, and its off-diagonal row sums are $\lambda_a \sum_j A_{aj} = \lambda_a < \theta_a + \lambda_a$, so it is strictly diagonally dominant with positive diagonal. A strictly diagonally dominant Z-matrix with positive diagonal is a non-singular M-matrix, hence $M^{-1} \geq 0$ entrywise [7]; existence and uniqueness follow, and $\vec{c} = M^{-1} \mathbf{1} > 0$ because no row of M^{-1} can vanish. Define

$$\frac{1}{\theta_M} = \sum_a c_a, \quad \bar{\theta} = m \theta_M, \quad (8)$$

and call $\bar{\theta}$ the effective market coefficient. Market clearing $\sum_a \vec{z}_a = \vec{v}$ then gives $\vec{\mu} - R_f \vec{p} = \theta_M V \vec{v}$, that is

$$\vec{p} = \frac{\vec{\mu} - \theta_M V \vec{v}}{R_f}. \quad (9)$$

Equation (9) closes the model. Prices are explicit in the primitives $(\vec{\mu}, V, \vec{v}, R_f)$ and in the single scalar θ_M , and no fixed point has to be solved, because θ_M is determined by (7) and (8) from $(\vec{\theta}, \vec{\lambda}, A)$ alone and does not depend on prices. This is a property of CARA preferences with Gaussian payoffs, not an approximation: demands are independent of initial wealth, so the wealth distribution never feeds back. Prices are positive, and the equilibrium admissible, exactly when

$$\vec{\mu} > \theta_M V \vec{v} \quad \text{componentwise}, \quad (10)$$

which is the only parameter restriction the closure imposes; parameters satisfying it are called admissible. If in addition $\vec{\mu} \gg 0$, every sufficiently small θ_M is admissible. All the results below concern θ_M , so they are statements about (9).

3.2 Returns and the price of risk

Gross returns are $R_i = X_i/p_i$. Writing $\mathbb{D}_p = \text{diag}(\vec{p})$, the return covariance and the vector of expected excess returns are

$$\Sigma = \mathbb{D}_p^{-1} V \mathbb{D}_p^{-1}, \quad \vec{r}_e = \mathbb{D}_p^{-1} (\vec{\mu} - R_f \vec{p}), \quad (11)$$

both endogenous. Let $W_M = \vec{p}^\top \vec{v}$ be the market value of the risky assets and $\vec{w} = \mathbb{D}_p \vec{v} / W_M$ its vector of value weights. Substituting (9) into (11) gives

$$\vec{r}_e = \theta_M W_M \Sigma \vec{w}, \quad (12)$$

the mean–covariance relation, now with every object in it determined. The market portfolio has excess return $R_M - R_f$ with $\mathbb{E}[R_M] - R_f = \theta_M \bar{\nu}^\top V \bar{\nu} / W_M$ and $\text{Var}(R_M) = \bar{\nu}^\top V \bar{\nu} / W_M^2$, so

$$\frac{\mathbb{E}[R_M] - R_f}{\text{Var}(R_M)} = \theta_M W_M, \quad S_M = \frac{\mathbb{E}[R_M] - R_f}{\text{sd}(R_M)} = \frac{\bar{\theta}}{m} \sqrt{\bar{\nu}^\top V \bar{\nu}}. \quad (13)$$

The first is the market price of risk per unit of variance; the second is the Sharpe ratio, and it is $\bar{\theta}$ multiplied by a constant that lives entirely on the payoff side. Every comparative static in $\bar{\theta}$ below is therefore a comparative static in the price of risk, in the strict sense. The unconditional statements are the aggregate ones: a fall in $\bar{\theta}$ lowers the market Sharpe ratio and the market risk premium. Statements about individual prices are conditional, since $\partial p_i / \partial \theta_M = -(V \bar{\nu})_i / R_f$, so that a fall in $\bar{\theta}$ raises p_i exactly when $(V \bar{\nu})_i > 0$ — when asset i covaries positively with the market payoff. Assets that hedge the market move the other way.

Remark 1 (On the riskless rate). Taking R_f as given amounts to a riskless asset in perfectly elastic supply, and then $W_M = \bar{p}^\top \bar{\nu}$ is an endogenous market value that need not equal aggregate initial wealth. If instead the riskless asset is in zero net supply, the budget identities give $\sum_a (W_a^0 - \bar{p}^\top z_a) = \sum_a W_a^0 - W_M$, so clearing that market requires $W_M = \sum_a W_a^0$ and determines

$$R_f = \frac{\bar{\mu}^\top \bar{\nu} - \theta_M \bar{\nu}^\top V \bar{\nu}}{\sum_a W_a^0}.$$

Both closures leave (9)–(13) intact, since neither touches θ_M ; they differ only in whether W_M or R_f is the free quantity. Under the second, the market value of the risky assets equals aggregate initial wealth, and since $\partial R_f / \partial \theta_M = -\bar{\nu}^\top V \bar{\nu} / \sum_a W_a^0 < 0$, comparison *raises* the riskless rate whenever it lowers $\bar{\theta}$. That closure also fixes $\bar{\nu}^\top \bar{p} = \sum_a W_a^0$, so prices cannot all move in the same direction: a fall in $\bar{\theta}$ redistributes value across assets instead of raising the aggregate.

4 Neutrality under common net risk aversion

Theorem 1. *Let Assumptions 1 and 2 hold and suppose $\theta_a = \theta$ for every a , that is, $\alpha_a(1 - \gamma_a) = \theta$. Then $\vec{c} = \theta^{-1} \mathbf{1}$ is the unique solution of (7), so every investor holds $\bar{z}_a = V^{-1}(\bar{\mu} - R_f \bar{p}) / \theta$ and $\bar{\theta} = \theta$, for every influence matrix satisfying Assumption 1 and every profile $(\lambda_1, \dots, \lambda_m)$.*

Proof. Substituting $\vec{c} = \theta^{-1} \mathbf{1}$ in row a of (7) and using $\sum_j A_{aj} = 1$ gives $(\theta + \lambda_a) / \theta - \lambda_a / \theta = 1$. Uniqueness follows from strict diagonal dominance, as above. $\square$

The hypothesis is that the product $\alpha_a(1 - \gamma_a)$ be common, so neutrality holds on the whole level set $\{(\alpha, \gamma) : \alpha(1 - \gamma) = \theta\}$. Along that set an investor with a stronger comparison motive is also more risk averse in own wealth, in exactly the proportion that

cancels. That is the content of the theorem, and it is weaker than it sounds if θ_a is read as the conventional coefficient of absolute risk aversion, which by Definition 1 it is not.

The cancellation occurs row by row, before any aggregation, which is why nothing about A or about the dispersion of λ_a enters. Economically, the benchmark is fixed in a unilateral deviation: under Assumption 2 the positions of others do not respond. Neutrality arises instead because at the equilibrium profile the benchmark carries the same risky exposure as the investor himself, so that $(\theta + \lambda_a)\bar{z} - \lambda_a A\bar{z} = (\theta + \lambda_a)\bar{z} - \lambda_a \bar{z} = \theta\bar{z}$. The benchmark exposure offsets the additional λ_a loading on own equilibrium wealth in the level of total risky exposure, not through any conjectured response to the deviation.

Theorem 1 is the network extension of the CARA case of [12], and we claim no more for it than that. His Proposition 2 identifies common cautiousness — here, common θ — as what makes the weighting drop out; the addition is that the weighting may be agent-specific and the comparison intensities may differ, and neither disturbs the conclusion.

Remark 2 (Comparison at constant own-wealth risk aversion). Theorem 1 varies λ_a holding θ fixed. The other natural path holds α fixed and varies γ_a , so that $\theta_a = \alpha(1 - \gamma_a)$ is dispersed whenever γ_a is. Neutrality then fails, but not through a new mechanism: the demand scales solve (7) with $\theta_a = \alpha(1 - \gamma_a)$ and $\lambda_a = \alpha\gamma_a$, and dispersion of the keeping-up weight is dispersion of net risk aversion. The two paths are therefore two coordinates of one statement. What neutrality requires is that $\alpha_a(1 - \gamma_a)$ be common; what breaks it is dispersion in that product, however generated.

Gómez's Proposition 2 requires more, in a second way. It asks that the preference parameters be identical across agents and, barring uniform weights, that the equilibrium value of the endowment be identical as well. Initial wealth is absent from (6) altogether, since CARA demands have no wealth effects, so no restriction on the endowment distribution is needed here.

Remark 3 (Two notions of neutrality). Theorem 1 asserts that individual positions are unchanged, which is stronger than saying the effective coefficient is unchanged. The latter requires only that $\sum_a c_a$ be unchanged, and deviations can offset. The distinction is not hypothetical: it occurs inside the example of Proposition 2, whose $\bar{\theta}(\lambda)$ returns to its no-comparison value at $\lambda = \frac{1}{3}$. There

$$\bar{c}(0) = (2, 1, \frac{1}{2}), \quad \bar{c}(\frac{1}{3}) = (\frac{5}{3}, \frac{7}{6}, \frac{2}{3}), \quad \bar{\theta}(0) = \bar{\theta}(\frac{1}{3}) = \frac{6}{7},$$

so the aggregate is exactly neutral while individual positions move by -16.7 , $+16.7$ and $+33.3$ per cent. Results below concern $\bar{\theta}$, so the weaker notion is the relevant one, and statements about it should not be read as statements about allocations.

5 Heterogeneous net risk aversion

We now let θ_a vary, which is exactly the configuration in which the aggregation conditions of [12] fail. The system (7) still has a unique positive solution, but it is no longer proportional to $\mathbf{1}$, and $\bar{\theta}$ depends on λ and on A .

Throughout this section $\lambda_a = \lambda$ is common, so that λ indexes a one-parameter family of economies running from no comparison to complete comparison. Two limits of that family can be computed in closed form for *every* network, and between them they organise everything that follows. Write

$$\vec{t} = \mathbb{D}_\theta^{-1} \mathbf{1}, \quad t_a = 1/\theta_a, \quad (14)$$

for the vector of risk tolerances, so that $\vec{c}(0) = \vec{t}$ and $\bar{\theta}(0) = H(\theta)$ is the harmonic mean of net risk aversion.

5.1 The marginal effect of comparison

Theorem 2. *Let Assumption 1 hold and $\lambda_a = \lambda$. Then*

$$\bar{\theta}'(0) = \frac{m \vec{t}^\top (\mathbb{I} - A) \vec{t}}{(\mathbf{1}^\top \vec{t})^2}. \quad (15)$$

If in addition A is doubly stochastic — not necessarily symmetric — then

$$\vec{t}^\top (\mathbb{I} - A) \vec{t} = \frac{1}{2} \sum_{a,j} A_{aj} (t_a - t_j)^2 \geq 0, \quad (16)$$

so $\bar{\theta}'(0) \geq 0$, with equality if and only if $t_a = t_j$ whenever $A_{aj} > 0$. For A merely row stochastic the sign of (15) is unrestricted.

Proof. Write $M(\lambda) = \mathbb{D}_\theta + \lambda(\mathbb{I} - A)$, so that $\vec{c} = M^{-1} \mathbf{1}$ and $\vec{c}' = -M^{-1}(\mathbb{I} - A)\vec{c}$. At $\lambda = 0$ we have $M(0) = \mathbb{D}_\theta$ and $\vec{c}(0) = \vec{t}$, hence $\mathbf{1}^\top \vec{c}'(0) = -\mathbf{1}^\top \mathbb{D}_\theta^{-1}(\mathbb{I} - A)\vec{t} = -\vec{t}^\top (\mathbb{I} - A)\vec{t}$, using $\mathbf{1}^\top \mathbb{D}_\theta^{-1} = \vec{t}^\top$. Since $\bar{\theta} = m/(\mathbf{1}^\top \vec{c})$, $\bar{\theta}' = -m(\mathbf{1}^\top \vec{c}')/(\mathbf{1}^\top \vec{c})^2$, which gives (15). For the second claim, row and column sums of A both equal one, so $\sum_{a,j} A_{aj} (t_a^2 + t_j^2)/2 = \frac{1}{2}(\sum_a t_a^2 + \sum_j t_j^2) = \vec{t}^\top \vec{t}$, and therefore $\frac{1}{2} \sum_{a,j} A_{aj} (t_a - t_j)^2 = \vec{t}^\top \vec{t} - \vec{t}^\top A \vec{t}$, which is (16). For the last claim take

$$A = \begin{pmatrix} 0 & 1 & 0 \\ 1 & 0 & 0 \\ 0 & 1 & 0 \end{pmatrix}, \quad \vec{\theta} = (1, 1, 2), \quad \vec{t} = (1, 1, \tfrac{1}{2}),$$

which satisfies Assumption 1 and has column sums $(1, 2, 0)$. Here $A\vec{t} = \mathbf{1}$, so $\vec{t}^\top A\vec{t} = \frac{5}{2}$ while $\vec{t}^\top \vec{t} = \frac{9}{4}$, giving $\vec{t}^\top (\mathbb{I} - A)\vec{t} = -\frac{1}{4}$ and $\bar{\theta}'(0) = 3(-\frac{1}{4})/(\frac{5}{2})^2 = -\frac{3}{25} < 0$. $\square$

In that example the two least risk averse investors absorb all the attention and the most risk averse receives none, so $\vec{\pi} = (\frac{1}{2}, \frac{1}{2}, 0)$ and $\vec{\pi}^\top \vec{\theta} = 1 < \frac{6}{5} = H(\theta)$. The effect is negative at the origin and the limit lies below the starting value, and $\bar{\theta}$ in fact falls monotonically from $6/5$ to 1. It is the complement of Proposition 2 below, where the marginal effect is positive and the decline sets in only at higher intensity.

Equation (16) is a Dirichlet form: the marginal effect of comparison is a weighted sum of squared differences in risk tolerance across the links of the network. It is positive

exactly when investors look at somebody unlike themselves, and the relevant notion of “unlike” is difference in tolerance, not in aversion. Double stochasticity is what makes the quadratic form a Dirichlet form; economically it says that attention is balanced, each investor receiving as much of it as he gives. It is the balance of attention, not the profile of tolerances, that the condition constrains: for a given $\vec{t}$ the form may well be non-negative without it, but double stochasticity is what makes it non-negative for *every* $\vec{t}$.

5.2 The limit of intense comparison

The other end of the family is governed by the stationary distribution of A read as a Markov chain. Let $C_1, \dots, C_r$ be its recurrent classes, $\vec{\pi}^{(k)}$ the stationary distribution on C_k normalised to sum to one, and $h_k(a)$ the probability that a chain started at a is absorbed into C_k .

Lemma 1. $c_a \leq 1/\min_b \theta_b$ for every a and every $\lambda \geq 0$.

Proof. Let $c^* = c_{a^*}$ be the largest entry. Row a^* of (7) reads $(\theta_{a^*} + \lambda)c^* = 1 + \lambda \sum_j A_{a^*j} c_j \leq 1 + \lambda c^*$, so $\theta_{a^*} c^* \leq 1$. $\square$

Theorem 3. Let Assumption 1 hold and $\lambda_a = \lambda$. Then

$$\lim_{\lambda \rightarrow \infty} \bar{\theta}(\lambda) = \frac{m}{\sum_{k=1}^r \frac{\mathbf{1}^\top \vec{h}_k}{\vec{\pi}^{(k)\top} \vec{\theta}}}. \quad (17)$$

In particular, if A has a single recurrent class with stationary distribution $\vec{\pi}$, then $\bar{\theta}(\lambda) \rightarrow \vec{\pi}^\top \vec{\theta}$.

Proof. The kernel of $\mathbb{I} - A$ is spanned by the absorption vectors $\vec{h}_1, \dots, \vec{h}_r$, which satisfy $\sum_k \vec{h}_k = \mathbf{1}$. Dividing $M(\lambda)\vec{c} = \mathbf{1}$ by λ gives $(\lambda^{-1}\mathbb{D}_\theta + \mathbb{I} - A)\vec{c} = \lambda^{-1}\mathbf{1}$, and $\vec{c}$ is bounded by Lemma 1, so any limit point $\vec{c}_\infty$ satisfies $(\mathbb{I} - A)\vec{c}_\infty = 0$, that is $\vec{c}_\infty = \sum_k \kappa_k \vec{h}_k$. To identify κ_k , extend $\vec{\pi}^{(k)}$ by zero off C_k and left-multiply $M(\lambda)\vec{c} = \mathbf{1}$ by $\vec{\pi}^{(k)\top}$. Since C_k is closed, $\vec{\pi}^{(k)\top} A = \vec{\pi}^{(k)\top}$, so the term in λ vanishes identically and $\vec{\pi}^{(k)\top} \mathbb{D}_\theta \vec{c} = 1$ for every λ . Passing to the limit and using $h_l = \delta_{kl}$ on C_k gives $\kappa_k \vec{\pi}^{(k)\top} \vec{\theta} = 1$. Then $\mathbf{1}^\top \vec{c}_\infty = \sum_k \kappa_k \mathbf{1}^\top \vec{h}_k$, and $\bar{\theta} = m/(\mathbf{1}^\top \vec{c})$ gives (17). The limit point is unique, so the whole family converges. $\square$

Theorem 3 is the organising result of the paper. Under intense comparison the aggregate is no longer the harmonic mean of net risk aversion but a *stationary* average: influence, not population, does the weighting. Reading $\bar{\theta}(0) = H(\theta)$ against (17) gives at once a sufficient condition for the comparison to lower the effective coefficient,

$$\vec{\pi}^\top \vec{\theta} < H(\theta) \implies \bar{\theta}(\lambda) < \bar{\theta}(0) \text{ for all large } \lambda, \quad (18)$$

in the single-class case. Reversal is therefore not an artefact of small examples but a statement about the correlation between net risk aversion and stationary influence: it

occurs when the investors who attract attention in the long run are less risk averse than the population as a whole. The three results that follow are the three ways (17) can turn out.

5.3 The mean-field limit

Proposition 1. *Let $\lambda_a = \lambda$ for all a and consider a sequence of economies indexed by m in which the empirical distribution of θ_a converges weakly to a law. Suppose the networks satisfy*

$$\theta_a \geq \underline{\theta} > 0 \text{ for all } a, m, \quad \delta_m = \max_a \left\| A_{a\cdot} - \frac{1}{m} \mathbf{1} \right\|_1 \longrightarrow 0. \quad (19)$$

Then, writing $\bar{\theta}_m$ for the effective coefficient of the m -investor economy,

$$\bar{\theta}_m(\lambda) \longrightarrow \bar{\theta}_\infty(\lambda) = H(\theta + \lambda) - \lambda, \quad H(x) = (\mathbb{E}[1/x])^{-1}. \quad (20)$$

The limit $\bar{\theta}_\infty$ is differentiable wherever the moments below exist, and

$$\frac{d\bar{\theta}_\infty}{d\lambda} = \frac{\text{Var}[(\theta + \lambda)^{-1}]}{\mathbb{E}[(\theta + \lambda)^{-1}]^2} \geq 0, \quad (21)$$

with equality if and only if θ is degenerate. If moreover $\mathbb{E}[\theta^3] < \infty$ then $\bar{\theta}_\infty = \mathbb{E}[\theta] - \text{Var}(\theta)/\lambda + O(\lambda^{-2})$ as $\lambda \rightarrow \infty$. Equation (21) is the derivative of the limit; we do not claim it is the limit of the derivatives of the finite economies.

Proof. Since $\theta_a \geq \underline{\theta}$, the variable $(\theta + \lambda)^{-1}$ is bounded by $(\underline{\theta} + \lambda)^{-1}$, so weak convergence alone gives convergence of $\mathbb{E}[(\theta + \lambda)^{-1}]$ and no separate integrability hypothesis is needed. By Lemma 1, $\|\bar{c}\|_\infty \leq 1/\underline{\theta}$ uniformly in m and λ . Hence $|\sum_j A_{aj}c_j - \bar{c}_m| = |(A_{a\cdot} - m^{-1}\mathbf{1})^\top \bar{c}| \leq \delta_m \|\bar{c}\|_\infty \leq \delta_m/\underline{\theta} \rightarrow 0$ uniformly in a , which is the mixing condition in endogenous form and is where the uniform bound is needed. Under it, row a of (7) becomes $(\theta_a + \lambda)c_a = 1 + \lambda\bar{c}$ in the limit, so $c_a = (1 + \lambda\bar{c})/(\theta_a + \lambda)$. Averaging and writing $M = \mathbb{E}[(\theta + \lambda)^{-1}]$ gives $\bar{c} = M/(1 - \lambda M)$ and $\bar{\theta} = 1/M - \lambda$. Differentiating gives (21), whose sign is Cauchy–Schwarz, with equality only for degenerate θ . The expansion follows from $\mathbb{E}[(\theta + \lambda)^{-1}] = \lambda^{-1}(1 - \mathbb{E}[\theta]/\lambda + \mathbb{E}[\theta^2]/\lambda^2 + O(\lambda^{-3}))$, which requires the third moment to control the remainder. $\square$

In this limit comparison interpolates the aggregation of risk aversion between the two classical aggregators: the harmonic mean at $\lambda = 0$, which is the standard CARA result, and the arithmetic mean as $\lambda \rightarrow \infty$. Since $H(\theta) \leq \mathbb{E}[\theta]$, the largest increase in the effective coefficient available is the harmonic–arithmetic gap of the net risk-aversion distribution, a statistic of the cross-section rather than a free parameter.

Condition (19) is primitive: it constrains only A and the support of θ . The leave-one-out network $A_{aj} = 1/(m-1)$, $j \neq a$, has $\delta_m = 2/m$ exactly, so it qualifies; exact equality at finite m would require identical rows, which Assumption 1 forbids since $A_{aa} = 0$. More generally δ_m is twice the total-variation distance of each row from the uniform

distribution, so (19) asks that every investor's benchmark be close to the population average — which is what “mean field” should mean, and is checkable from the network alone. The uniform lower bound $\underline{\theta}$ cannot be dropped: it is what converts closeness of rows into closeness of benchmarks, since $\vec{c}$ is controlled only through Lemma 1.

Remark 4 (Consistency with Theorem 3). Condition (19) forces the stationary distribution towards the uniform one, and (17) then returns $\vec{\pi}^\top \vec{\theta} \rightarrow \mathbb{E}[\theta]$, which is the $\lambda \rightarrow \infty$ end of (20). The mean-field regime is thus the case of Theorem 3 in which influence is spread evenly, and $\bar{\theta}'(0) \geq 0$ in it because near-uniform rows are near doubly stochastic, so (16) applies approximately. Both directions of the interpolation are consequences of the two limits.

5.4 Non-monotonicity and sign reversal

Proposition 2. *There exist normalised benchmarks and net risk-aversion profiles for which $d\bar{\theta}/d\lambda < 0$ on a non-empty interval and for which $\bar{\theta}(\lambda) < \bar{\theta}(0)$ at sufficiently large comparison intensities.*

Proof. Take $m = 3$, $\theta = (\frac{1}{2}, 1, 2)$, $\lambda_a = \lambda$ and

$$A = \begin{pmatrix} 0 & 1 & 0 \\ 1 & 0 & 0 \\ 1 & 0 & 0 \end{pmatrix},$$

which satisfies Assumption 1. Solving (7) gives

$$\bar{\theta}(\lambda) = \frac{3(\lambda + 2)(3\lambda + 1)}{12\lambda^2 + 24\lambda + 7}, \quad \bar{\theta}'(\lambda) = -\frac{3(12\lambda^2 + 6\lambda - 1)}{(12\lambda^2 + 24\lambda + 7)^2}. \quad (22)$$

Hence $\bar{\theta}'(0) = 3/49 > 0$ and $\bar{\theta}'(\lambda) < 0$ for $\lambda > (\sqrt{21} - 3)/12 \approx 0.132$. The function rises from $\bar{\theta}(0) = 6/7 \approx 0.857$ to a maximum ≈ 0.860 at that point, then falls, returning to $6/7$ at $\lambda = 1/3$ and continuing downwards: $\bar{\theta}(1) = 0.837209$ and $\bar{\theta}(2) = 0.815534$. $\square$

Both features of the example are read off the two limits of Section 5.1 and Section 5.2, and neither depends on the closed form. The column sums of A are $(2, 1, 0)$, so A is not doubly stochastic and (16) does not apply; (15) nonetheless returns $\bar{\theta}'(0) = 3/49 > 0$, so comparison initially raises the effective coefficient. At the other end, investor 3 is transient and $\{1, 2\}$ is the unique recurrent class, with $\vec{\pi} = (\frac{1}{2}, \frac{1}{2}, 0)$. Theorem 3 therefore gives

$$\lim_{\lambda \rightarrow \infty} \bar{\theta}(\lambda) = \vec{\pi}^\top \vec{\theta} = \frac{1}{2}(\frac{1}{2} + 1) = \frac{3}{4} < \frac{6}{7} = H(\theta) = \bar{\theta}(0),$$

so (18) holds and reversal is forced. The economics is now explicit. The most risk averse investor watches the least risk averse and is watched by nobody, so stationary influence assigns him weight zero: his risk aversion is present in the harmonic mean at $\lambda = 0$ and absent from the stationary mean at $\lambda = \infty$. Non-monotonicity is the collision of the two

limits — a positive marginal effect at the origin with a strictly lower value at infinity — and the closed form only locates the turning point.

How often reversal occurs is a question about a distribution over networks, and the answer is a property of that distribution rather than of the mechanism. Under a dense i.i.d.-weight design, specified in the supplement, it becomes rare as the number of investors grows. That ensemble does not satisfy (19): normalised i.i.d. rows stay a fixed distance from the uniform row, since $\sum_k U_{ak} \sim m/2$ makes $\delta_m \rightarrow \mathbb{E}|2U - 1| = \frac{1}{2}$. Incoming influence and stationary weights nonetheless become increasingly balanced, and the numerical mixing deviation $\max_a |\sum_j A_{aj}c_j - \bar{c}|$ falls with m while $\bar{\theta}'(0)$ approaches the mean-field value $\text{CV}^2(\vec{t})$; what does this is not convergence of rows but a law of large numbers, the weights being drawn independently of the types. We make no formal convergence claim for the ensemble. Nothing in this licenses a claim about networks in general — least of all because the independence it relies on is exactly what homophily destroys. Real influence structures carry communities, hubs and correlation between characteristics and position, and Remark 5 shows that reversal survives at any population size once such structure is present. A supplement reports the frequency of both the cumulative event $\bar{\theta}(0.5) < \bar{\theta}(0)$ and the marginal event $\bar{\theta}'(0.5) < 0$ over 20,000 draws at each of ten combinations of population size and dispersion, together with the generators, the seeds and the analytic derivative used. Reversal is already below one per cent at $m = 6$ and indistinguishable from zero at $m \geq 10$ in that ensemble.

Remark 5 (Reversal at any population size). The decay reported in the supplement reflects mixing, not size. Because (7) is local, a block-diagonal A decouples: each block solves its own system and $\bar{\theta} = (m^{-1} \sum_a c_a)^{-1}$ is the reciprocal of the population-average capacity. Replicating an identical block leaves that average, and hence $\bar{\theta}$, unchanged. Taking k disjoint copies of the network of Proposition 2 therefore gives $\bar{\theta}(\lambda)$ independent of k , so that $\bar{\theta}(1) = 0.837209 < 6/7 = \bar{\theta}(0)$ at $m = 3k$ for every k ; we verify this numerically up to $m = 3000$. The construction is robust to weak ties across blocks: mixing each row with a fraction ϵ of the uniform complete network preserves the reversal at $m = 30$ for $\epsilon \leq 0.1$ ($\bar{\theta}(1) = 0.8526$ at $\epsilon = 0.1$) and destroys it by $\epsilon = 0.2$ ($\bar{\theta}(1) = 0.8681$). The simulation should therefore be read as describing a dense ensemble in which local benchmarks approach a common average, and not as a claim about large networks in general: block structure, or any persistent correlation between type and network position, sustains the reversal at arbitrary population size.

Note the contrast with Proposition 3, which uses the same block structure to the opposite end. There net risk aversion is constant *within* blocks and the mechanism vanishes; here it is heterogeneous within blocks and absent *across* them, and the mechanism survives at any scale. What separates the two is where the heterogeneity sits relative to the partition, not the topology.

5.5 Perfect homophily blocks the mechanism

Proposition 3. *Partition the investors into groups such that $A_{aj} = 0$ whenever a and j belong to different groups, and suppose θ_a is constant within each group. Then $c_a = 1/\theta_a$*

solves (7) for every profile $(\lambda_1, \dots, \lambda_m)$, and $\bar{\theta}$ equals the harmonic mean of θ regardless of the intensity of comparison.

Proof. Row a of (7) reads $(\theta_a + \lambda_a)c_a - \lambda_a \sum_j A_{aj}c_j = 1$. If $c_j = 1/\theta_j$ and $\theta_j = \theta_a$ for every j with $A_{aj} > 0$, the left-hand side is $(\theta_a + \lambda_a)/\theta_a - \lambda_a/\theta_a = 1$. Uniqueness gives the claim, and $\bar{\theta} = m/\sum_a(1/\theta_a)$ is the harmonic mean of θ . $\square$

Theorem 3 explains why. Each group is a recurrent class, so $r > 1$, and θ is constant on C_k , whence $\vec{\pi}^{(k)\top} \vec{\theta} = \theta_k$ whatever the stationary distribution inside the class. With every investor in some class, $\mathbf{1}^\top \vec{h}_k$ is the size m_k of the class and (17) reads $m/\sum_k m_k/\theta_k = H(\theta)$. The two limits therefore coincide, which is the strongest possible form of neutrality of the aggregate: not that $\bar{\theta}$ returns to $H(\theta)$, but that it never leaves. With two equally sized groups at $\theta \in \{0.5, 2\}$, $\bar{\theta} = 0.8$ at $\lambda = 0$, at $\lambda = 1$ and at $\lambda = 1000$; the arithmetic mean 1.25 that Proposition 1 would predict is never approached.

The hypothesis is exact segregation together with exact constancy of θ within blocks, and both are needed: if θ varies inside a class then $\vec{\pi}^{(k)\top} \vec{\theta} \neq \theta_k$ and (17) moves, as Remark 5 exploits. So the correct statement is that *perfect homophily by net risk-aversion type* suppresses the mechanism, not that homophily does. It is worth knowing because assortativity on wealth and on financial sophistication is a documented feature of social networks [15, 6].

6 Conclusion

Relative wealth concerns in a CARA economy aggregate through the influence network in a way that two limits organise. The marginal effect at zero comparison is a quadratic form in risk tolerances, which is a Dirichlet form — and so non-negative — exactly when attention is balanced. The limit under intense comparison is a stationary average of individual coefficients, weighted by influence rather than by population. They anchor the two ends of the family rather than tracing the path between them, but between them they account for the mean-field interpolation from the harmonic to the arithmetic mean, for non-monotonicity and sign reversal, and for the complete inertness of the aggregate under perfect homophily by type. What decides which case obtains is the correlation between net risk aversion and stationary influence.

Because the economy is closed, these are statements about prices and not only about a conditional relation. Equilibrium prices are $\vec{p} = (\vec{\mu} - \theta_M V \vec{v})/R_f$ and the market Sharpe ratio is $\bar{\theta} \sqrt{\vec{v}^\top V \vec{v}}/m$, both explicit, because CARA demands carry no wealth effects and θ_M therefore does not depend on prices. A network that lowers $\bar{\theta}$ compresses the market risk premium and the market Sharpe ratio. Its effect on an individual price has the sign of $(V \vec{v})_i$, so assets that covary positively with the market payoff rise while hedges fall; and under the zero-net-supply closure, where $\vec{v}^\top \vec{p}$ is pinned to aggregate initial wealth, the riskless rate rises and value is redistributed across assets rather than created.

One boundary should be kept in view. Here θ_a is risk aversion net of the comparison, not the coefficient of absolute risk aversion in own wealth, so “common risk aversion”

in Theorem 1 means a common value of the product $\alpha_a(1 - \gamma_a)$, along which stronger comparison motives come paired with greater own-wealth risk aversion.

The practical reading is that a representative-agent treatment of relative concerns reproduces the mean-field case and is safe only where influence is close to evenly spread. Where attention is unbalanced — where some investors are watched much more than they watch — it can err in sign. Two questions follow naturally. The first is quantitative: how far can $\vec{\pi}^\top \vec{\theta}$ fall below $H(\theta)$ for networks with realistic degree distributions and realistic correlation between risk aversion and centrality? The second is empirical: stationary influence is in principle measurable, and Theorem 3 converts it into a prediction about the aggregate.

Replication

All numerical results are reproducible from the repository

<https://github.com/renatovicente/price-of-risk-aggregation>

The module `scripts/network.py` solves (7); `scripts/neutrality.py` verifies Theorem 1 and Proposition 1; `scripts/sign_reversal.py` produces the table in the supplement, including the network and type generators and the seeds; `scripts/block_reversal.py` produces the block replication and the mixing robustness of Remark 5; and `scripts/characterization.py` checks Theorems 2 and 3, the Dirichlet identity (16), the bound of Lemma 1 and each of the three cases of (17); and `scripts/closed_economy.py` checks the first-order conditions (6) against direct numerical optimisation, market clearing, the price formula (9), the recovery of (12) and the Sharpe identity (13). The scripts depend only on `numpy` and `scipy`, read no data and draw from fixed seeds; `./reproduce.sh` runs all six in about fifteen seconds.

References

- [1] A. B. Abel, *Asset prices under habit formation and catching up with the Joneses*, American Economic Review Papers and Proceedings 80 (1990) 38–42.
- [2] C. Ballester, A. Calvó-Armengol, Y. Zenou, *Who’s who in networks. Wanted: the key player*, Econometrica 74 (2006) 1403–1417.
- [3] J. Y. Campbell, J. H. Cochrane, *By force of habit: a consumption-based explanation of aggregate stock market behavior*, Journal of Political Economy 107 (1999) 205–251.
- [4] Y. L. Chan, L. Kogan, *Catching up with the Joneses: heterogeneous preferences and the dynamics of asset prices*, Journal of Political Economy 110 (2002) 1255–1285.

- [5] L. De Gennaro Aquino, E. G. De Giorgi, Y. Lou, M. S. Strub, *Investment and asset pricing with relative wealth concerns and multiple risky assets*, The North American Journal of Economics and Finance 84 (2026) 102628.
- [6] M. Bailey, R. Cao, T. Kuchler, J. Stroebe, *The economic effects of social networks: evidence from the housing market*, Journal of Political Economy 126 (2018) 2224–2276.
- [7] A. Berman, R. J. Plemmons, *Nonnegative Matrices in the Mathematical Sciences*, SIAM, Philadelphia, 1994.
- [8] P. DeMarzo, R. Kaniel, I. Kremer, *Diversification as a public good: community effects in portfolio choice*, Journal of Finance 59 (2004) 1677–1716.
- [9] B. Dupor, W.-F. Liu, *Jealousy and equilibrium overconsumption*, American Economic Review 93 (2003) 423–428.
- [10] J. Galí, *Keeping up with the Joneses: consumption externalities, portfolio choice, and asset prices*, Journal of Money, Credit and Banking 26 (1994) 1–8.
- [11] C. Gollier, *Wealth inequality and asset pricing*, Review of Economic Studies 68 (2001) 181–203.
- [12] J.-P. Gómez, *The impact of keeping up with the Joneses behavior on asset prices and portfolio choice*, Finance Research Letters 4 (2007) 95–103.
- [13] H. Kraft, A. Meyer-Wehmann, F. T. Seifried, *Dynamic asset allocation with relative wealth concerns in incomplete markets*, Journal of Economic Dynamics and Control 113 (2020) 103857.
- [14] D. Lacker, T. Zariwopoulou, *Mean field and n -agent games for optimal investment under relative performance criteria*, Mathematical Finance 29 (2019) 1003–1038.
- [15] M. McPherson, L. Smith-Lovin, J. M. Cook, *Birds of a feather: homophily in social networks*, Annual Review of Sociology 27 (2001) 415–444.
- [16] S. Panageas, *The implications of heterogeneity and inequality for asset pricing*, NBER Working Paper 26974 (2020).

Supplementary Material

Relative Wealth Concerns, Network Aggregation, and Effective Risk Aversion

Renato Vicente

This supplement reports the simulation referred to in Section 5.4 of the paper. Notation follows the paper throughout: $\bar{\theta} = m / \sum_a c_a$ is the effective market coefficient, $\vec{c}$ solves $[\mathbb{D}_{\theta+\lambda} - \mathbb{D}_\lambda A] \vec{c} = \mathbf{1}$ with $\lambda_a = \lambda$ common, θ_a is net risk aversion and A the row-stochastic influence matrix.

S1 How often does reversal occur?

Proposition 2 of the paper exhibits a network on which comparison eventually lowers $\bar{\theta}$, and Theorem 3 gives the condition $\bar{\pi}^\top \vec{\theta} < H(\theta)$ under which this happens. Neither says how common such networks are. That is a question about a distribution over networks rather than about the mechanism, and the answer depends entirely on which distribution is chosen. We report one.

Design

Fix $\lambda = 0.5$ and draw 20,000 independent economies at each of ten combinations of population size m and dispersion σ . Off-diagonal entries of A are uniform on $(0, 1)$ with zero diagonal, normalised by rows; net risk aversion is $\theta_a = \max\{1 + \sigma Z_a, 0.15\}$ with Z_a standard normal. The same generative rule is used at every m , and the population size is held fixed within each row; the resulting size pattern is therefore a property of this sequence of dense random-network ensembles rather than of population size as such. Seeds are $1000m + 10\sigma$, set in `scripts/sign_reversal.py`.

Two distinct events are recorded and should not be conflated:

$$\text{cumulative: } \bar{\theta}(0.5) < \bar{\theta}(0), \quad \text{marginal: } \bar{\theta}'(0.5) < 0. \quad (\text{S1})$$

The derivative is computed analytically rather than by finite differences. Writing $M(\lambda) = \mathbb{D}_\theta + \lambda(\mathbb{I} - A)$ and $\vec{c} = M^{-1}\mathbf{1}$,

$$\vec{c}' = -M^{-1}(\mathbb{I} - A)\vec{c}, \quad \bar{\theta}' = -\frac{m \mathbf{1}^\top \vec{c}'}{(\mathbf{1}^\top \vec{c})^2}, \quad (\text{S2})$$

so each draw costs two linear solves with the same factorisation of M . Intervals are Wilson intervals at 95 per cent.

m	σ	cumulative: $\theta(0.5) < \theta(0)$		marginal: $\theta'(0.5) < 0$	
		frequency	95% interval	frequency	95% interval
3	0.4	0.0919	[0.0879, 0.0959]	0.0952	[0.0912, 0.0993]
3	0.8	0.0340	[0.0316, 0.0366]	0.0372	[0.0346, 0.0399]
6	0.4	0.0057	[0.0047, 0.0068]	0.0062	[0.0052, 0.0073]
6	0.8	0.0008	[0.0005, 0.0014]	0.0008	[0.0005, 0.0012]
10	0.4	0.0001	[0.0000, 0.0004]	0.0001	[0.0000, 0.0004]
10	0.8	0.0000	[0.0000, 0.0002]	0.0000	[0.0000, 0.0002]
20	0.4	0.0000	[0.0000, 0.0002]	0.0000	[0.0000, 0.0002]
20	0.8	0.0000	[0.0000, 0.0002]	0.0000	[0.0000, 0.0002]
60	0.4	0.0000	[0.0000, 0.0002]	0.0000	[0.0000, 0.0002]
60	0.8	0.0000	[0.0000, 0.0002]	0.0000	[0.0000, 0.0002]

Table S1: Frequency of negative cumulative and marginal effects at $\lambda = 0.5$, over 20,000 draws per row with Wilson intervals. All ten cells of the design are reported; entries of 0.0000 are exact zeros, no draw out of 20,000 producing the event, and 0.0002 is the upper endpoint of the 95% Wilson interval when no event is observed.

Results

Under this design reversal is already below one per cent at $m = 6$ and indistinguishable from zero at $m \geq 10$. The two events move together, which is what one expects when $\bar{\theta}$ is close to monotone over $[0, 0.5]$, and the cumulative frequency is slightly the smaller of the two at every cell.

What the table does and does not show

The decay in m is a property of this ensemble. The ensemble does *not* satisfy the primitive mixing condition of Proposition 1 of the paper. With $A_{aj} = U_{aj} / \sum_{k \neq a} U_{ak}$ and U_{aj} uniform, the law of large numbers gives $\sum_{k \neq a} U_{ak} \sim m/2$, so

$$\left\| A_{a\cdot} - \frac{1}{m} \mathbf{1}^\top \right\|_1 \approx \frac{1}{m} \sum_j |2U_{aj} - 1| \rightarrow \mathbb{E} |2U - 1| = \frac{1}{2}, \quad (\text{S3})$$

and $\delta_m \not\rightarrow 0$. The decay in the table therefore cannot be deduced from that proposition, and we claim no formal convergence for this ensemble. What we do observe is that the endogenous deviation $\max_a |\sum_j A_{aj} c_j - \bar{c}|$ falls steadily with m — from 0.14 at $m = 10$ to 0.011 at $m = 4000$ — and that $\bar{\theta}'(0)$ approaches the mean-field value $\text{CV}^2(\vec{t})$ to three decimals by $m = 4000$. The reason is a law of large numbers rather than any convergence of rows: the weights are drawn independently of the types, so $\sum_j A_{aj} t_j \rightarrow \bar{t}$ even though $A_{a\cdot}$ stays far from uniform. Making the weights depend on types breaks it at once — with weights proportional to t_j^3 the same deviation is 3.99 at $m = 4000$ against 0.03 under independence — which is the sense in which this ensemble is the opposite of a homophilous one. The decay is a property of that independence, not of size. Remark 5

of the paper makes this precise: replicating the network of Proposition 2 in disjoint blocks leaves $\bar{\theta}(\lambda)$ unchanged for every population size, verified up to $m = 3000$, and the reversal survives weak ties across blocks up to a mixing fraction of about $\epsilon = 0.1$. Real influence structures carry communities, hubs and correlation between characteristics and network position, none of which this ensemble contains.

The two dispersions are not a clean comparative static. Reversal is rarer at $\sigma = 0.8$ than at $\sigma = 0.4$, but the difference cannot be read as an effect of dispersion alone. The truncation at 0.15 binds for 14.4 per cent of draws at $\sigma = 0.8$ against 1.7 per cent at $\sigma = 0.4$, which changes the shape of the distribution of θ_a and raises $\mathbb{E}[\theta]$ from 1.00 to 1.06. We report both rows and draw no conclusion from their difference.

Nothing here bears on the theorems. Theorems 2 of the paper hold for every network satisfying Assumption 1. The table describes the measure of networks on which the two limits point in opposite directions under one particular sampling scheme, which is a question of prevalence and not of validity.

Replication

The table is produced by `scripts/sign_reversal.py` and the block construction of Remark 5 by `scripts/block_reversal.py`, in the repository

<https://github.com/renatovicente/price-of-risk-aggregation>

Both are self-contained, depend only on `numpy` and `scipy`, read no data and draw from fixed seeds; `./reproduce.sh` runs them together with the rest in about fifteen seconds.